

Jason (Jia) Teoh

jiateoh at gmail | linkedin.com/in/jiateoh | https://jiateoh.github.io
Los Angeles, CA

Education

Ph.D., Computer Science

9/16 - Present

Advisor: Dr. Miryung Kim

UC Los Angeles, GPA: 3.940 (ongoing)

M.S., Computer Science

9/16 - 3/19

UC Los Angeles, GPA: 3.933

B.A., Computer Science

8/10 – 5/13

UC Berkeley, GPA: 3.868 (High Distinction)

Research Experience

Software Engineering and Analysis Laboratory (SEAL)

3/18 - Present

- Investigate precision and performance debugging of big data systems such as Apache Spark with a focus on data impact.
- Combine latency instrumentation with data provenance to produce record-level performance provenance which associates application performance with input data and program definitions.
- Utilize user-defined function semantics to improve data provenance trace precision in identifying relevant inputs that produce suspicious or faulty output records in dataflow applications.

Scalable Analytics Institute (ScAi)

9/16 - 3/18

- Investigated mobile dataflow computing and bridging the gap between mobile and cloud computing
- Developed Android framework integrated with Spark Catalyst Optimizer to automatically export sub-plans and integrate with Spark jobs through Apache Kafka.

Publications

[SoCC 2020] Jason Teoh, Muhammad Ali Gulzar, and Miryung Kim. 2020. Influence-Based Provenance for Dataflow Applications with Taint Propagation. *ACM Symposium on Cloud Computing (SoCC '20)*. 24.4% acceptance rate

[SoCC 2019] Jason Teoh, Muhammad Ali Gulzar, Guoqing Harry Xu, and Miryung Kim. 2019. PerfDebug: Performance Debugging of Computation Skew in Dataflow Systems. *ACM Symposium on Cloud Computing (SoCC '19)*. 24.8% acceptance rate

Project Experience

Cloud Spanner (<https://cloud.google.com/spanner/>)

6/19-9/19

Software Engineering Intern - Google

- Migrated an internal diagnostics infrastructure from JavaScript to TypeScript.
- Designed and implemented an automated rule-based diagnostics framework to analyze query executions and identify performance problems or anomalies. Detected issues are presented to users of the tool alongside additional documentation and potential fixes.
- Collaborated with Spanner developers to identify dozens of query execution characteristics that typically indicate poor performance. Implemented and refined a small subset of these rules for demonstration.
- Integrated the rule-based analysis framework with the existing diagnostics infrastructure to visually bind detected issues with corresponding query execution information.

Thirdeye (github.com/linkedin/pinot/tree/master/thirdeye)

10/15 - 9/16

Senior Software Engineer - LinkedIn

- Open source anomaly detection framework for monitoring business metrics.
 - Utilizes Pinot (realtime distributed OLAP datastore with horizontal scaling and low latency).
- Designed and implemented multifaceted dashboard for straightforward anomaly investigation.
- Investigated and designed solutions to improve anomaly detection scalability.
- Onboarded new use cases from data bootstrapping to server setup.
- Start[in]: Coached new college graduates in workplace topics such as communication and owning one's role.

LinkedIn Segmentation and Targeting Tool (LISTT)

6/13 - 9/15

Senior Software Engineer (5/15- 9/15), Software Engineer (6/13 - 5/15) - LinkedIn

- Suite of self-service applications for targeting audiences and computing user-defined member-level attributes.
- Enhanced computational flow to improve data quality and application robustness for over 500 metrics per each of LinkedIn's 400 million members, on a daily or weekly refresh cycle.
- Heterogenous data integration for hundreds of data definitions defined by SQL (Teradata), Hive, Pig.
- Leveraged Apache Datafu Hourglass to reduce redundant computation and optimize up data delivery.
- Independently developed flow-monitoring CLI and graph UI for time and stability analysis of application
- Mentored intern and onboarded new hires.
- Attended Hadoop Summit (2014-2015) to learn about big data technologies such as Spark, Mesos, YARN.

RaINier

5/12 - 8/12

Software Engineer Intern - LinkedIn

- Developed self-service business analytics dashboard tool for viewing key metrics and dimensional splits.
- Implemented multidimensional (OLAP) cube caching layer with background refresh daemons.
- Presented project to a variety of audiences: project fair, business users, and prospective developers.

Teaching

Software Engineering

Teaching Assistant (CS 130) - UCLA

Fall '20

- Lead class discussions (Zoom) and online (Piazza), held office hours, administered grades, supervised term-long group projects from proposal to final presentation.

Database Systems

Teaching Assistant (CS 143) - UCLA

Fall '17

- Lead discussions in class and online (Piazza), held office hours, administered grades, managed class projects, conducted extra study reviews.

Selected Coursework

Distributed Systems, Databases, Machine Learning, Software Engineering, Operating Systems, Artificial Intelligence, Statistical Programming, Probability, Efficient Algorithms and Intractable Problems, Compilers

Honors/Awards

Member of Upsilon Pi Epsilon (Berkeley chapter)

Spring '12

Dean's Honors List (UC Berkeley)

Fall '11 to Spring '13

Skills

Programming: Scala, Java, TypeScript, Python, SQL, Bash, Apache Pig, Apache Hive, Javascript, Ruby, R, C, LaTeX, Scheme, TI-Basic

Frameworks/Libraries: Spark, Hadoop, Kafka, Spring, Lucene, AngularJS, jQuery, Sigma (js), Datafu Hourglass, Ruby on Rails, Maven, Gradle